

**Congratulations! You passed!**

Grade received 100% To pass 80% or higher

[Go to next item](#)**Final Exam**

Latest Submission Grade 100%

1.

1 / 1 point

**Question: Modify a table**Modify the columns in table `sales` to match the provided schema

Given the schema:

column	type
datetime	BIGINT
disc_code	INT
product_code	STRING
unit_cost	STRING
unit_price	STRING

Select the SQL query that achieves the following:

- Selects columns `datetime`, `unit_price`, and `disc_code`
- Converts `datetime` column to `TIMESTAMP`; `rename as datetime`
- Converts `unit_price` to `FLOAT`
- Writes results to a table named `discounts`



```
1 --OPTION
2 DROP TABLE IF EXISTS discounts;
3 CREATE TABLE discounts AS
4 SELECT
5   to_timestamp(datetime) as datetime,
6   disc_code,
7   CAST(unit_price AS FLOAT)
8 FROM
9   sales;
10 SELECT
11   *
12 FROM
13   discounts
```



```
1 --OPTION
2 CREATE
3 OR REPLACE TEMPORARY VIEW discounts AS
4 SELECT
5   to_date(datetime) as datetime,
6   disc_code,
7   CAST(unit_price AS FLOAT)
8 FROM
9   sales;
10 SELECT
11   *
12 FROM
13   discounts
```



```
1 --OPTION
2 CREATE
3 OR REPLACE TEMPORARY VIEW discounts AS
4 SELECT
5   to_timestamp(time) as datetime,
6   disc_code,
7   unit_price AS FLOAT
8 FROM
9   sales;
10 SELECT
11   *
12 FROM
```

```
13    discounts
```

```
1  --OPTION
2  DROP TABLE IF EXISTS discounts;
3  CREATE TABLE discounts AS
4  SELECT
5      to_date(datetime),
6      disc_code,
7      unit_price AS FLOAT
8  FROM
9      sales;
10 SELECT
11     *
12 FROM
13     discounts;
```

Correct

2.

1 / 1 point

### Question: Basic Math

Create a new column for the table, `customers`, by applying basic math operations to an existing column.

Given the schema:

column	type
cust_first_name	STRING
cust_last_name	STRING
cust_city	STRING
cust_state	STRING
cust_country_id	STRING
cust_income	BIGINT
cust_credit_lim	BIGINT
overdue	STRING

Select the SQL query that achieves the following:

- Includes columns: `cust_first_name`, `cust_last_name`, `cust_city`, `cust_state`, `overdue`, and `overdue_charge`
- Converts the `overdue` column to `BOOLEAN`
- Creates a new column `overdue_charge` which should show 2% of the `cust_credit_lim` column, rounded to the nearest cent
- Stores results to a temporary view named `overdue_payments`

```
1  --OPTION
2  CREATE OR REPLACE TEMPORARY VIEW overdue_payments AS
3  SELECT cust_first_name,
4      cust_last_name,
5      cust_city,
6      cust_state,
7      CAST(overdue AS BOOLEAN),
8      ROUND(cust_credit_lim * .02, 2) AS overdue_charge
9  FROM customers;
10
11  SELECT * FROM overdue_payments;
```

```
1  --OPTION
2  CREATE OR REPLACE TEMPORARY VIEW overdue_payments AS
3  SELECT cust_first_name,
4      cust_last_name,
5      cust_city,
6      cust_state,
7      overdue AS BOOLEAN,
8      ROUND(cust_credit_lim * .02) AS overdue_charge
9  FROM customers;
10
11  SELECT * FROM overdue_payments;
```

```

1  --OPTION
2  CREATE OR REPLACE TEMPORARY VIEW overdue_payments AS
3  SELECT cust_first_name,
4    cust_last_name,
5    cust_city,
6    cust_state,
7    CAST(overdue AS BOOLEAN),
8    cust_credit_lim * .02
9  FROM customers;
10
11  SELECT * FROM overdue_payments;

```

```

1  --OPTION
2  CREATE OR REPLACE TEMPORARY VIEW overdue_payments AS
3  SELECT cust_first_name,
4    cust_last_name,
5    cust_city,
6    cust_state,
7    CAST(overdue) BOOLEAN,
8    cust_credit_lim * .02
9  FROM customers;
10
11  SELECT * FROM overdue_payments;

```

Correct

3.

1 / 1 point

## Question: Define a table with schema

Write a table, `restaurant_reviews`, and define its schema in the creation statement.

Select the table creation statement that achieves the following:

- Uses the path: `/mnt/training/restaurants/reviews.csv`
- Allows the first line to be read as a header
- Specifies that that table is created using `csv`
- Creates a table with the following schema:

column	type
restaurant_id	STRING
user_id	INT
rating	FLOAT
text_review	STRING
time_recorded	INT

```

1  --OPTION
2  DROP TABLE IF EXISTS restaurant_reviews;
3  CREATE TABLE restaurant_reviews (
4    restaurant_id,
5    user_id,
6    rating,
7    text_review,
8    time_recorded
9  ) USING csv OPTIONS (
10    path "mnt/training/restaurants/reviews.csv",
11  );

```

1 --OPTION

```

1 --CREATE
2 DROP TABLE IF EXISTS restaurant_reviews;
3 CREATE TABLE restaurant_reviews (
4   restaurant_id STRING,
5   user_id INT,
6   rating FLOAT,
7   text_review STRING,
8   time_recorded INT
9 ) USING csv OPTIONS (
10   path "/mnt/training/restaurants/reviews.csv",
11   header "true"
12 );

```



```

1 --OPTION
2 CREATE OR REPLACE TEMPORARY VIEW restaurant_reviews (
3   restaurant_id STRING,
4   user_id INT,
5   rating FLOAT,
6   text_review STRING,
7   time_recorded INT
8 ) USING csv OPTIONS (
9   path "/mnt/training/restaurants/reviews.csv",
10  header "true"
11 );

```



```

1 --OPTION
2 CREATE OR REPLACE TEMPORARY VIEW restaurant_reviews (
3   restaurant_id,
4   user_id,
5   rating,
6   text_review,
7   time_recorded
8 ) USING parquet OPTIONS (
9   path "mnt/training/restaurants/reviews.csv",
10  header "true"
11 );
12

```

Correct

4.

1 / 1 point

## Question: Remove duplicates and sort

Deduplicate and sort restaurant review information from the table, `restaurant_evals`

Given the following schema:

column	type
restaurant_id	STRING
user_id	INT
rating	DOUBLE
text_review	STRING
time_recorded	BIGINT

Select the query that achieves the following:

- Removes duplicate rows
- Casts `time_recorded` as a timestamp. This value can be converted from unixtime
- Sorts rows by `rating` in descending order, and then by `time_recorded` in ascending order
- Displays **only** columns `restaurant_id`, `user_id`, `rating`, `text_review`, `time`



```

1 --OPTION
2 SELECT
3   DISTINCT *
4   FROM
5   (
6     SELECT
7       restaurant_id,
8       user_id,
9       rating

```

```
    text_review,
11   from_unixtime(time_recorded) as time
12   FROM
13   restaurant_evals
14 )
15 ORDER BY
16   rating DESC,
17   time ASC
```



```
--OPTION
1 SELECT
2   DISTINCT *
3   FROM
4   (
5     SELECT
6       restaurant_id,
7       user_id,
8       rating,
9       text_review,
10      from_unixtime(time_recorded) as time
11     FROM
12     restaurant_evals
13   )
14 ORDER BY
15   rating
16   time
```



```
--OPTION
1 SELECT
2   DISTINCT *,
3   from_unixtime(time_recorded) as time
4   FROM
5   restaurant_evals
6 ORDER BY
7   rating DESC,
8   time ASC
```



```
--OPTION
1 SELECT
2   DISTINCT *,
3   from_unixtime(time_recorded) as time
4   FROM
5   restaurant_evals
6 ORDER BY
7   rating,
8   time
```

 Correct

5.

1 / 1 point

## Question: Limit Results

Return the top 10 results for data that matches the given criteria.

Given the schema:

column	type
employee_id	STRING
time_recorded	TIMESTAMP
net_gain	DOUBLE

Write a SQL query on the table `q1_sales` that achieves the following:

- Includes the columns: `employee_id`, `closing_date`, and `net_gain`
- Casts `time_recorded` column as date and renames to `closing_date`
- Sorts rows by the `net_gain` column in descending order
- Limits the results to the top 10 `net_gain`

○

```
1 --OPTION
2 SELECT
3   employee_id,
4   to_date(time_recorded) closing_date,
5   net_gain
6 FROM
7   q1_sales
8 ORDER BY
9   net_gain DESC
10
```

○

```
1 --OPTION
2 SELECT
3   employee_id,
4   to_date(time_recorded),
5   net_gain
6 FROM
7   q1_sales
8 ORDER BY
9   net_gain DESC
10 LIMIT
11   10
```

○

```
1 --OPTION
2 SELECT
3   *
4 FROM
5   q1_sales
6 ORDER BY
7   net_gain DESC
8 LIMIT
9   10
```

○

```
1 --OPTION
2 SELECT
3   employee_id,
4   to_date(time_recorded) closing_date,
5   net_gain
6 FROM
7   q1_sales
8 ORDER BY
9   net_gain DESC
10 LIMIT
11   10
```

Correct

6.

1 / 1 point

## Question: Join Tables

Perform a join on two tables `games` and `goals`.

Given the schemas:

`orders`

column	type
order_id	INT
order_date	DATE
cust_id	STRING

`customers`

column	type
cust_id	STRING
cust_name	STRING
contact_name	STRING
country	STRING

Select the query that achieves the following:

- Performs an inner join on `orders` and `customers` on the column `cust_id`
- Includes **only** the following columns: `cust_id`, `order_date`, and `cust_name`



```

1 --OPTION
2 SELECT
3   orders.cust_id,
4   order_date,
5   cust_name
6 FROM
7   orders
8   JOIN customers ON orders.cust_id = customers.cust_id

```



```

1 --OPTION
2 SELECT
3   orders.cust_id,
4   order_date,
5   cust_name
6 FROM
7   orders
8   JOIN customers ON cust_id = customers.cust_id

```



```

1 --OPTION
2 SELECT
3   cust_id,
4   order_date,
5   cust_name
6 FROM
7   orders
8   JOIN customers ON cust_id = cust_id

```



```

1 --OPTION
2 SELECT
3   cust_id,
4   order_date,
5   cust_name
6 FROM
7   orders
8   JOIN customers ON cust_id = customers.cust_id

```

**Correct**

7.

1 / 1 point

## Question: Union Tables

Perform a union of two tables `q1_earnings` and `q2_earnings`

Assume that these two tables have the same schema.

Select the query that achieves the following:

- Includes all records from both tables, including any duplicates



```
1 --OPTION
2 SELECT *
3 FROM q1_earnings
4 UNION JOIN
5 SELECT *
6 FROM q2_earnings
```



```
1 --OPTION
2 SELECT DISTINCT *
3 FROM q1_earnings
4 UNION ALL
5 SELECT *
6 FROM q2_earnings
```



```
1 --OPTION
2 SELECT *
3 FROM q1_earnings
4 UNION ALL
5 SELECT *
6 FROM q2_earnings
```



```
1 --OPTION
2 SELECT *
3 FROM q1_earnings
4 UNION
5 SELECT *
6 FROM q2_earnings
```

Correct

8.

1 / 1 point

### Question: Extract year and month

Extract the year and month from the `timestamp` field in the table `timetable1`

Given the schema:

column	type
time	timestamp
rating	DOUBLE

Select the SQL query that achieves the following:

- Extracts the year and month from the timestamp
- Includes columns `year`, `month`, and `rating` in the results



```
1 --OPTION
2 SELECT
3   month(time) month,
4   day(time) day,
5   rating
6 FROM
7   | timetable1
```

```
1 --OPTION
2 SELECT
3   CAST(time AS month) month,
4   CAST(time AS day) day,
5   rating
6 FROM
7   timetable1
```

```
1 --OPTION
2
3 SELECT time "MM",
4 time "dd",
5 rating
6 FROM timetable1
```

```
1 --OPTION
2 SELECT month(time),
3 day(time),
4 rating
5 FROM timetable1
```

 Correct

9,

1 / 1 point

### Question: Extract day of week

Create a new `day` column from the given table, `sales`.

Given the schema:

column	type
time	TIMESTAMP
purchase_amt	DOUBLE
sale_id	INT

Select the query that achieves the following:

- Creates a new column, `day`, that contains the day of the week associated with the timestamp.  
This column should contain the name of the day, e.g. Sunday, Monday, Tuesday.
- Returns three columns: `day`, `purchase_amt`, and `sale_id`

```
1 --OPTION
2 SELECT
3   date_format(time, "E"),
4   purchase_amt,
5   sale_id
6 FROM
7   sales
```

```
1 --OPTION
2 SELECT
3   day(to_date(time), "E") day,
4   purchase_amt,
5   sale_id
6 FROM
7   sales
```



```
1 --OPTION
2 SELECT
3   date_format(to_date(time), "E") day,
4   purchase_amt,
5   sale_id
6 FROM
7   sales
```

**Correct**



```
1 --OPTION
2 SELECT
3   date_format(time, "E") day,
4   purchase_amt,
5   sale_id
6 FROM
7   sales
```

**Correct**

10.

1 / 1 point

### Question: Minimum Function

Display the minimum `total_rev` for each `category` in the table `feb_sales`.

Given the schema:

column	type
item_id	INT
category	STRING
sale_price	DOUBLE
quantity	INT

Select the query that achieves the following:

- Creates a new column `total_revenue` by multiplying `sale_price` by the `quantity`, rounded to the nearest cent.
- Displays the minimum `total_revenue` in each `category`. Label that column `min_revenue`
- Displays the sum of all `total_revenues` in each `category`. Label that column `category_sum`
- Includes columns: `category`, `category_sum`, `min_revenue`



```
1 --OPTION
2 WITH total_rev AS (
3   SELECT
4     `,
5     ROUND(sale_price * quantity, 2) total_revenue
6   FROM
7     feb_sales
8 )
9 SELECT
10   category,
11   ROUND(SUM(total_revenue), 2) category_sum,
12   MIN(total_revenue) min_revenue
13   FROM
14   total_rev
15   GROUP BY
16   category
```

**Correct**



```
1 --OPTION
2 SELECT
3   -----
```

```

3   category,
4   ROUND(SUM(total_revenue), 2) category_sum,
5   MIN(total_revenue) min_revenue
6   FROM
7   (
8     SELECT
9       *,
10      ROUND(sale_price * quantity, 2) total_revenue
11     FROM
12     feb_sales
13   )
14 GROUP BY
15   category,
16   total_revenue

```



```

1 --OPTION
2 SELECT
3   category,
4   ROUND(SUM(total_revenue), 2) category_sum,
5   MIN(total_revenue) min_revenue
6   FROM (
7     SELECT *
8     ROUND(sale_price * quantity, 2) total_revenue
9     FROM feb_sales
10   )
11 GROUP BY category
12

```

Correct



```

1 --OPTION
2 SELECT
3   category,
4   ROUND(SUM(total_revenue), 2) category_sum,
5   MIN(total_revenue) min_revenue
6   FROM
7   (
8     SELECT
9       *
10      ROUND(sale_price * quantity, 2) total_revenue
11     FROM
12     feb_sales
13   )

```

### 11. Question: Maximum Function

1/1 point

Display the maximum `total_revenue` for each `category` and `item_id` in the table `feb_sales`.

Given the schema:

column	type
item_id	INT
category	STRING
sale_price	DOUBLE
quantity	INT

Select the query that achieves the following:

- Creates a new column `total_revenue` by multiplying `sale_price` by the `quantity`, rounded to the nearest cent.
- Display the maximum `total_revenue` for each `category` and `item_id`. Label that column `max_revenue`.



```

1 --OPTION
2 SELECT
3   category,
4   item_id,
5   ROUND(SUM(total_revenue), 2) category_sum,
6   MAX(total_revenue) max_revenue
7   FROM (
8     SELECT *,
9     ROUND(sale_price * quantity, 2) total_revenue
10    FROM feb_sales
11   )
12 GROUP BY category, item_id

```



```
1 --OPTION
2 SELECT
3   category,
4   item_id,
5   ROUND(SUM(total_revenue), 2) category_sum,
6   MAX(total_revenue) max_revenue
7   FROM (
8     SELECT *,
9       ROUND(sale_price * quantity, 2) total_revenue
10    FROM feb_sales
11  )
12 GROUP BY category
```



```
1 --OPTION
2 SELECT
3   category,
4   item_id,
5   ROUND(SUM(total_revenue), 2) category_sum,
6   MAX(total_revenue) max_revenue
7   FROM (
8     SELECT *,
9       ROUND(sale_price * quantity, 2) total_revenue
10    FROM feb_sales
11  )
12
```



```
1 --OPTION
2 WITH total_rev AS (
3   SELECT
4     *,
5     ROUND(sale_price * quantity, 2) total_revenue
6   FROM
7     feb_sales
8 )
9 SELECT
10   category,
11   ROUND(SUM(total_revenue), 2) category_sum,
12   MAX(total_revenue) max_revenue
13   FROM
14   total_rev
15 GROUP BY
16   category item_id
```

Correct

12.

1 / 1 point

## Question: Average Function

Display the average `total_revenue` for each `category` in the table `feb_sales`.

Given the schema:

column	type
item_id	INT
category	STRING
sale_price	DOUBLE
quantity	INT

Select the query that achieves the following:

- Creates a new column `total_revenue` by multiplying `sale_price` by the `quantity`, rounded to the nearest cent.
- Displays the average `total_revenue` in each `category`. Label that column `avg_revenue`.
- Displays the standard deviation for `avg_revenue`. Label that column `std_deviation`.



```

1 --OPTION
2 SELECT
3   category,
4   ROUND(STD(total_revenue), 3) std_deviation,
5   ROUND(AVG(total_revenue),2) avg_revenue
6 FROM
7 (
8   SELECT
9     *
10    ROUND(sale_price * quantity, 2) total_revenue
11   FROM
12   feb_sales
13 )
14 GROUP BY
15   category,
16   item_id

```

```

1 --OPTION
2 SELECT
3   category,
4   ROUND(STD(total_revenue), 3) std_deviation,
5   ROUND(AVG(total_revenue),2) avg_revenue
6 FROM
7 (
8   SELECT
9     *,
10    ROUND(sale_price * quantity, 2) total_revenue
11   FROM
12   feb_sales
13 )
14 GROUP BY
15   category, quantity

```

```

1 --OPTION
2
3 SELECT
4   category,
5   ROUND(STD(total_revenue), 3) std_deviation,
6   ROUND(AVG(total_revenue),2) avg_revenue
7 FROM
8 (
9   SELECT
10    *,
11    ROUND(sale_price * quantity, 2) total_revenue
12   FROM
13   feb_sales
14 )
15

```

```

1 --OPTION
2 SELECT
3   category,
4   ROUND(STD(total_revenue), 3) std_deviation,
5   ROUND(AVG(total_revenue),2) avg_revenue
6 FROM
7 (
8   SELECT
9     *,
10    ROUND(sale_price * quantity, 2) total_revenue
11   FROM
12   feb_sales
13 )
14 GROUP BY
15   category

```

Correct

13.

1 / 1 point

### Question: Pivot table

Display categories as columns. Show total\_rev in each category.

Given the schema:

column	type
item_id	INT
category	STRING
sale_price	DOUBLE
quantity	INT

Select the query that achieves the following:

- Creates a new column `total_revenue` by multiplying `sale_price` by the `quantity`, rounded to the nearest cent.
- Display the sum of `total_revenue` in each `category`, where each `category` is a column.



```
1 --OPTION
2 SELECT
3 *
4 FROM
5 (
6   SELECT
7     category,
8     ROUND(SUM(total_revenue), 2)
9   FROM
10    feb_sales
11  ) PIVOT (
12    ROUND(sale_price * quantity, 2) total_revenue FOR category IN ('books', 'magazines', 'movies')
13 );
```



```
1 --OPTION
2 SELECT
3 *
4 FROM
5 (
6   SELECT
7     ROUND(sale_price * quantity, 2) total_revenue
8   FROM
9    feb_sales
10  ) PIVOT (
11    ROUND(SUM(total_revenue), 2) FOR category IN ('books', 'magazines', 'movies')
12 );
```



```
1 --OPTION
2 SELECT
3 *
4 FROM
5 (
6   SELECT
7     category,
8     ROUND(sale_price * quantity, 2) total_revenue
9   FROM
10    feb_sales
11  ) PIVOT (
12    ROUND(SUM(total_revenue), 2) FOR category IN ('books', 'magazines', 'movies')
13 );
```



```
1 --OPTION
2 SELECT
3 *
4 FROM
5 (
6   SELECT
7     *,
8     ROUND(sale_price * quantity, 2) total_revenue
9   FROM
10    feb_sales
11  ) PIVOT (
12    ROUND(SUM(total_revenue), 2) FOR category IN ('books', 'magazines', 'movies')
13 );
```

Correct

14.

1 / 1 point

-

## Question: Null Values and Aggregates

Compute sums of `sales_rev` grouped by `dept` after dropping null values from `inventory` table.

Given the schema:

column	type
item_id	STRING
dept	STRING
sales_rev	DOUBLE
price	DOUBLE

Select the query that achieves the following:

- Drops any rows that contain null values in either the `item_id` or the `dept` column
- Aggregates sums of the `sales_rev` column grouped by `dept`



```
1 --OPTION
2 SELECT
3   dept,
4   SUM(sales_rev)
5 FROM
6   inventory
7 GROUP BY
8   dept
```



```
1 --OPTION
2 SELECT
3   dept,
4   SUM(sales_rev)
5 FROM
6   inventory
7 WHERE
8   (
9     item_id IS NOT NULL
10    AND dept IS NOT NULL
11  )
12 GROUP BY
13   dept
```



```
1 --OPTION
2 SELECT
3   dept,
4   SUM(sales_rev)
5 FROM
6   inventory
7 WHERE
8   item_id AND dept IS NOT NULL
```



```
1 --OPTION
2 SELECT
3   dept,
4   SUM(sales_rev)
5 FROM
6   inventory
7 WHERE
8   item_id AND dept IS NOT NULL
9 GROUP BY
10  dept
```

Correct

15.

1 / 1 point

### Question: Generate subtotals by rollup

Compute averages of `revenue` grouped by `title` and `month` such that the results include averages across all months as well as a subtotal for an individual months from the `book_sales` table.

Given the schema:

column	type
title	STRING
month	INTEGER
revenue	DOUBLE

Select the query that achieves the following:

- Coalesces null values generated by the `ROLLUP` clause in the `month` column
- Creates a new column, `avg_revenue`, rounded to the nearest dollar
- Groups records by `title` and then `month`



```
1 --OPTION
2 SELECT
3   title,
4   month
5   ROUND(AVG(revenue), 2) as avg_revenue
6 FROM book_sales
7 GROUP BY ROLLUP (title, month)
```



```
1 --OPTION
2 SELECT
3   COALESCE(title, "All items") AS title,
4   COALESCE(month, "All months") AS month,
5   ROUND(AVG(revenue), 2) as avg_revenue
6 FROM book_sales
7 GROUP BY ROLLUP (title, month)
8 ORDER BY title, month;
9
```



```
1 --OPTION
2 SELECT
3   COALESCE(title AS title),
4   COALESCE(month AS month),
5   ROUND(AVG(revenue), 2) as avg_revenue
6 FROM book_sales
7 GROUP BY ROLLUP (title, month)
8 ORDER BY title, month;
9
10
```



```
1 --OPTION
2 SELECT
3   title,
4   month
5   ROUND(AVG(revenue), 2) as avg_revenue
6 FROM book_sales
7 GROUP BY ROLLUP (title, month)
8 ORDER BY title, month;
9
```

16.

1 / 1 point

Edit ✓ - ×

## Question: MapTypes

Flatten the `customers` table described below so that there is no `MAP` and the query results include 4 columns: `first_name`, `last_name`, `titles`, and `payment_methods`

Given the schema:

column	type
<code>first_name</code>	STRING
<code>last_name</code>	STRING
<code>subscriptions</code>	MAP<STRING, MAP <titles:ARRAY<STRING>, payment_methods:ARRAY<STRING> > >

Select the SQL query that achieves the following:

- Returns results with the following schema

column	type
<code>first_name</code>	STRING
<code>last_name</code>	STRING
<code>titles</code>	ARRAY
<code>payment_methods</code>	ARRAY

**Note:** Remember, when you `EXPLODE` a `MapType` column, it breaks the field into `key`, `value` pairs.



```

1 --OPTION
2 SELECT
3   first_name,
4   last_name,
5   EXPLODE (subscriptions),
6   key.titles,
7   key.payment_methods
8 FROM
9   customers

```



```

1 --OPTION
2 SELECT
3   first_name,
4   last_name,
5   subscriptions.titles,
6   subscriptions.payments
7 FROM
8   customers

```



```

1 --OPTION
2 WITH explode_subs AS (
3   SELECT
4     first_name,
5     last_name,
6     EXPLODE (subscriptions)
7   FROM
8     customers
9 )
10 SELECT
11   first_name,
12   last_name,
13   subscriptions.titles,
14   subscriptions.payment_methods
15 FROM
16   explode_subs

```



```

1 --OPTION
2 WITH explode_subs AS (
3   SELECT
4     first_name,
5     last_name,

```

```

6   EXPLODE (subscriptions)
7   FROM
8   customers
9 )
10 SELECT
11   first_name,
12   last_name,
13   value.titles,
14   value.payment_methods
15 FROM
16   explode_subs

```

Correct

17.

1 / 1 point

## Question: Arrays and Indexes

Access the first element of the `titles` array from the table `customers`.

Given the schema:

column	type
first_name	STRING
last_name	STRING
subscriptions	MAP<STRING, MAP <titles:ARRAY<STRING>, payment_methods:ARRAY<STRING> > >

Select the SQL query that achieves the following:

- Returns results that include only columns: `first_name`, `last_name`, and `title`
- Lists the first element of the `titles` array as the string in the `title` column



```

1 --OPTION
2 WITH explode_subs AS (
3   SELECT
4     first_name,
5     last_name,
6     EXPLODE (subscriptions)
7   FROM
8     customers
9 )
10 SELECT
11   first_name,
12   last_name,
13   subscriptions.titles[0]
14 FROM
15   explode_subs

```



```

1 --OPTION
2 WITH explode_subs AS (
3   SELECT
4     first_name,
5     last_name,
6     EXPLODE (subscriptions)
7   FROM
8     customers
9 )
10 SELECT
11   first_name,
12   last_name,
13   subscriptions.titles,
14   subscriptions.payment_methods
15 FROM
16   explode_subs

```



```

1 --OPTION
2 SELECT
3   first_name,
4   last_name,
5   value.titles[0] AS title
6 FROM (
7   |  |  | SELECT

```

```
8   first_name,  
9   last_name,  
10  EXPLODE (subscriptions)  
11  FROM  
12  customers  
13  )
```

```
1 --OPTION  
2 SELECT  
3   first_name,  
4   last_name,  
5   value.titles[0] AS title  
6 FROM  
7   SELECT  
8   first_name,  
9   last_name,  
10  EXPLODE (subscriptions)  
11  FROM  
12  customers  
13
```

Correct

18.

1 / 1 point

### Question: Transform

Transform mixed-case `titles` , in the table `books` , to upper case.

Given the schema:

column	type
titles	ARRAY

Each array contains mixed-case book titles, e.g. `[the red barn, Living, HAPPINESS]` .

Select the query that achieves the following:

- Returns the `titles` column with all values written in uppercase.

```
1 --OPTION  
2 SELECT  
3   TRANSFORM(t, titles -> UPPER(t)) titles  
4 FROM  
5   books
```

```
1 --OPTION  
2 SELECT  
3   TRANSFORM(titles, t -> UPPER(t)) titles  
4 FROM  
5   books
```

```
1 --OPTION  
2 SELECT  
3   TRANSFORM(titles, t -> UPPER(t))  
4 FROM  
5   books
```

```
1 --OPTION
2 SELECT
3   TRANSFORM(t, titles -> UPPER(t)) titles
```

Correct

19.

1 / 1 point

## Question: Flag Records

Create a new column to flag low battery levels with no available backups from the table `devices`

Given the schema:

column	type
battery_levels	ARRAY (INT)
backups	INT
location_id	INT

Select the query that achieves the following:

- Creates a new column named `needs_service` that holds Boolean values
- Uses a higher-order function to flag locations reporting battery levels under 4 **and** no available backups
- Returns columns `location_id` and `needs_service`



```
1 --OPTION
2 SELECT
3   location_id,
4   backups,
5   EXISTS(battery_levels, b -> b < 4 AND backups = 0) needs_service
6 FROM devices
```



```
1 --OPTION
2 SELECT
3   location_id,
4   backups,
5   FILTER(battery_levels, b -> b < 4 AND backups = 0)
6 FROM devices
```



```
1 --OPTION
2 SELECT
3   location_id,
4   FILTER(battery_levels, b -> b < 4 AND backups = 0) needs_service
5 FROM devices
```



```
1 --OPTION
2 SELECT
3   location_id,
4   EXISTS(battery_levels, b -> b < 4 AND backups = 0) needs_service
5 FROM devices
```

Correct

## Question: Reduce

Find the sum of the `battery_levels` in each array in the table `devices`

Given the schema:

column	type
battery_levels	ARRAY(INT)
backups	INT
location_id	INT

Select the query that achieves the following:

- Creates a new column, `total_batt_levels`, that holds the sum of all battery levels at a location
- Returns two columns: `location_id` and `total_batt_levels`



```
1 --OPTION
2 SELECT
3   location_id,
4   REDUCE(battery_levels, 0, (level, acc) -> level + acc) total_batt_levels
5 FROM
6   devices
```



```
1 --OPTION
2 SELECT
3   location_id,
4   REDUCE(level, 0, battery_levels -> level + acc) total_batt_levels
5 FROM
6   devices
```



```
1 --OPTION
2 SELECT
3   location_id,
4   REDUCE(battery_levels, 0, level, acc -> level + acc) total_batt_levels
5 FROM
6   devices
```



```
1 --OPTION
2 SELECT
3   location_id,
4   REDUCE(level, 0, (battery_levels, acc) -> level + acc) total_batt_levels
5 FROM
6   devices
```

 Correct